




Unit 3 - Process Deadlock

Introduction

- System consists of resources
- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - **request**
 - **use**
 - **release**


Introduction (contd.)

- A process in a multiprogramming system is said to be in dead lock if it is waiting for a particular event that will never occur
- Example:
 - All automobiles trying to cross
 - Traffic Completely stopped
 - Not possible without backing some



Resource Deadlock

- A process request a resource before using it, and release after using it
- If the resource is not available when it is requested, the requesting process is force to wait
- Process P1 holds resource R1 and needs resource R2 to continue; Process P2 holds resource R2 and needs resource R1 to continue – deadlock



Deadlock Characterization


- Deadlock can arise if four conditions hold simultaneously.
 - **Mutual exclusion:** only one process at a time can use a resource
 - **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes
 - **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task
 - **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Resource-Allocation Graph


- A set of vertices V and a set of edges E .
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system
- **request edge** – directed edge $P_i \rightarrow R_j$
- **assignment edge** – directed edge $R_j \rightarrow P_i$

Resource-Allocation Graph (Cont.)


➤ Process




➤ Resource Type with 4 instances




➤ P_i requests instance of R_j




➤ P_i is holding an instance of R_j



Example of a Resource Allocation Graph



Resource Allocation Graph With A Deadlock





Basic Facts

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock



Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state:
 - Deadlock prevention
 - Deadlock avoidance
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX



Deadlock Prevention

- **Mutual Exclusion** – not required for sharable resources (e.g., read-only files); must hold for non-sharable resources
 - (Making all resources shareable)
- **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 - Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none allocated to it.
 - Low resource utilization; starvation possible

Deadlock Prevention (contd.)

- **No Preemption** –

- If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
- Preempted resources are added to the list of resources for which the process is waiting
- Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

- **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration



Ostrich Algorithm

- Deadlock Ignorance Method
- A strategy of ignoring problem on the basis that they may be exceedingly rare
- “To stick one’s head in the sand and pretend there is no problem”
- Used when it is cost-effective to allow the problem rather than its prevention



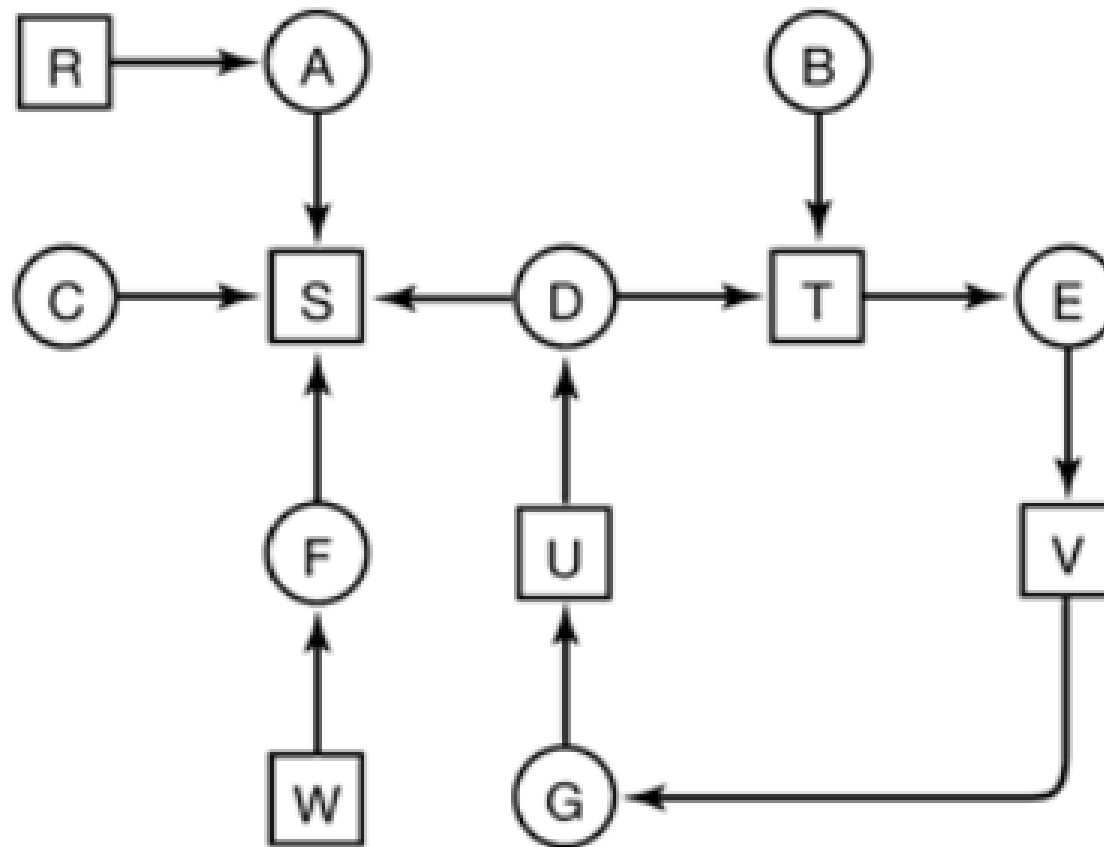
Deadlock Detection

- The system does not attempt to prevent deadlocks from occurring
- Instead, it lets them occur, tries to detect when this happens, and then takes some action to recover after the fact
- Different approaches:
 - Deadlock Detection with One Resource of Each Type
 - Deadlock Detection with Multiple Resource of Each Type

Deadlock Detection with One Resource of Each Type


- If resource graph consists one or more cycle then deadlock occurs
- If no cycle then no deadlock
- Here, one resource of each type exists (one scanner, one printer, etc.)
- Is this System Deadlocked? If so which process are involved?
 - Process A holds R and wants S
 - Process B holds nothing but wants T
 - Process C holds nothing but wants S
 - Process D holds U and wants S and T
 - Process E holds T and wants V
 - Process F holds W and wants S
 - Process G holds V and wants U

Deadlock Detection-One Resource (contd.)



Deadlock Detection-One Resource (contd.)

➤ After Extraction:




Deadlock Detection with Multiple Resources of Each Type

- Multiple copies of resources available
- A matrix-based algorithm in order to detect deadlock
- n processes. P_1 through P_n
- E is the existing resource vector
- A be the available resource vector
- two arrays, C , the current allocation matrix, and R , the request matrix
- C_{ij} is the number of instances of resource j that are held by process i
- R_{ij} is the number of instances of resource j that P_i wants

Deadlock Detection - Multiple Resources (contd.)

Resources in existence
($E_1, E_2, E_3, \dots, E_m$)

Current allocation matrix




| | | | | |
|----------|----------|----------|---------|----------|
| C_{11} | C_{12} | C_{13} | \dots | C_{1m} |
| C_{21} | C_{22} | C_{23} | \dots | C_{2m} |
| \vdots | \vdots | \vdots | | \vdots |
| C_{n1} | C_{n2} | C_{n3} | \dots | C_{nm} |

Row n is current allocation
to process n

Resources available
($A_1, A_2, A_3, \dots, A_m$)

Request matrix



| | | | | |
|----------|----------|----------|---------|----------|
| R_{11} | R_{12} | R_{13} | \dots | R_{1m} |
| R_{21} | R_{22} | R_{23} | \dots | R_{2m} |
| \vdots | \vdots | \vdots | | \vdots |
| R_{n1} | R_{n2} | R_{n3} | \dots | R_{nm} |

Row 2 is what process 2 needs

Deadlock Detection - Multiple Resources (contd.)

- The deadlock detection algorithm can now be given, as follows.
 - Look for an unmarked process, P_i , for which the i -th row of R is less than or equal to A
 - If such a process is found, add the i -th row of C to A , mark the process, and go back to step 1
 - If no such process exists, the algorithm terminates.

Example

| Process | Allocated | Request | Available |
|---------|-----------|---------|-----------------------|
| P0 | 0 1 0 | 0 0 0 | 0 0 0 |
| P1 | 2 0 0 | 2 0 2 | 0 1 0 |
| P2 | 3 0 3 | 0 0 0 | 3 1 3 |
| P3 | 2 1 1 | 1 0 0 | 5 2 4 |
| P4 | 0 0 2 | 0 0 2 | <u>5 2 6</u> 7 2 6 |

| A | B | C |
|---|---|---|
| 7 | 2 | 6 |

Work = available



Recovery from Deadlock

➤ Recovery through Preemption

- Temporarily take a resource away from its current owner and give it to another process

➤ Recovery through Rollback

- When a deadlock is detected, it is easy to see which resources are needed.
- . To do the recovery, a process that owns a needed resource is rolled back to a point in time before it acquired some other resource by starting one of its earlier checkpoints

➤ Recovery through Killing Processes

- Kill one or more processes

Deadlock Avoidance - Safe and Unsafe State

- A state is said to be safe if it is not deadlocked and there is a some scheduling order in which every process can run to completion
- In a safe state, system can guarantee that all processes will finish – no deadlock occur
- For an unsafe state, no such guarantee can be given – deadlock may occur.

| Has Max | | |
|---------|---|---|
| A | 3 | 9 |
| B | 2 | 4 |
| C | 2 | 7 |

Free: 3

(a)

| Has Max | | |
|---------|---|---|
| A | 3 | 9 |
| B | 4 | 4 |
| C | 2 | 7 |

Free: 1

(b)

| Has Max | | |
|---------|---|---|
| A | 3 | 9 |
| B | 0 | — |
| C | 2 | 7 |

Free: 5

(c)

| Has Max | | |
|---------|---|---|
| A | 3 | 9 |
| B | 0 | — |
| C | 7 | 7 |

Free: 0

(d)

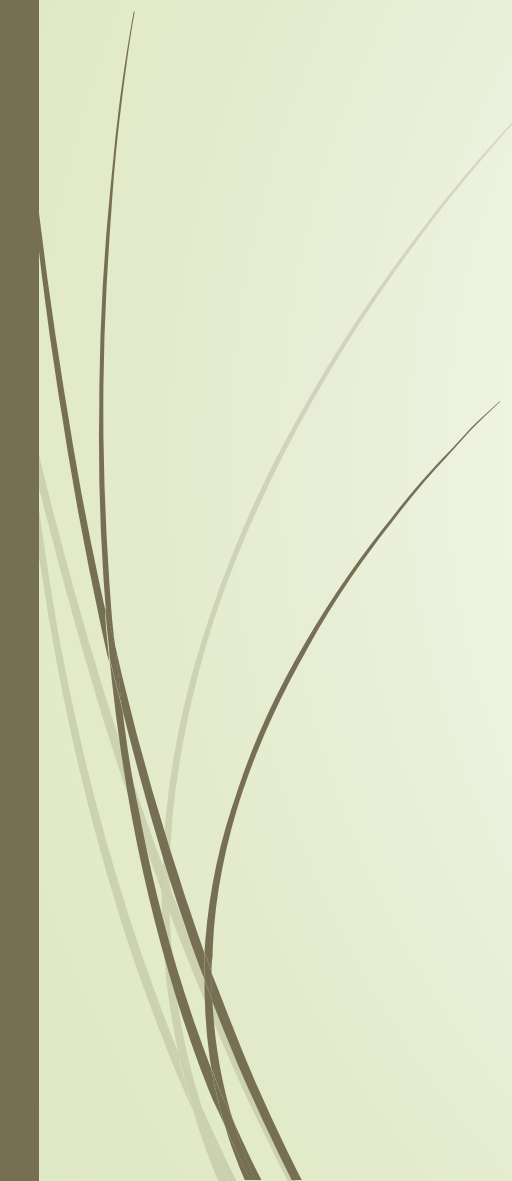
| Has Max | | |
|---------|---|---|
| A | 3 | 9 |
| B | 0 | — |
| C | 0 | — |

Free: 7

(e)



Banker's Algorithm

- This approach to the deadlock problem anticipates a deadlock before it actually occurs
 - This approach employs an algorithm to assess the possibility that deadlock could occur and act accordingly
 - It is named as Banker's algorithm because the process is analogous to that used by a banker in deciding if a loan can be safely made or not
- 

Banker's Algorithm (contd.)

- The following are the features that are to be considered for avoidance of the deadlocks per the Banker's Algorithms
- Each process declares maximum number of resources of each type that it may need
- Keep the system in a safe state in which we can allocate resources to each process in some order and avoid deadlock
- Check for the safe state by finding a safe sequence: where resources that P_i needs can be satisfied by available resources plus resources held by P_j where $j < i$
- Resource allocation graph algorithm uses claim edges to check for a safe state

Banker's Algorithm (contd.)

For Single Resource

- Consider an analogy in which 4 processes (P1, P2, P3 and P4) can be compared with the customers in a bank, resources such as printers etc. as cash available in the bank and the Operating system as the Banker

| Processes | Resources used | Maximum resources |
|-----------|----------------|-------------------|
| P1 | 0 | 6 |
| P2 | 0 | 5 |
| P3 | 0 | 4 |
| P4 | 0 | 7 |

| Processes | Resources used | Maximum resources |
|-----------|----------------|-------------------|
| P1 | 1 | 6 |
| P2 | 1 | 5 |
| P3 | 2 | 4 |
| P4 | 4 | 7 |

Banker's Algorithm (contd.)

For Multiple Resource

| Process | Allocated | Request | Available |
|---------|-----------|---------|----------------|
| P0 | 0 1 0 | 0 0 0 | 0 0 0 |
| P1 | 2 0 0 | 2 0 2 | 0 1 0 |
| P2 | 3 0 3 | 0 0 0 | 3 1 3 |
| P3 | 2 1 1 | 1 0 0 | <u>5 2 4</u> |
| P4 | 0 0 2 | 0 0 2 | 5 2 6 7 2 6 |

| A | B | C |
|---|---|---|
| 7 | 2 | 6 |

Work = available

Deadlock VS Starvation

| Deadlock | Starvation |
|---|--|
| 1) A deadlock is a condition in operating systems in which no process proceeds for execution and wait for resources that have been acquired by some other processes. It is also known as circular wait. | 1) Starvation is a situation when a process keeps waiting (and starving) for a resource that is being held by other high priority processes. |
| 2) In a deadlock, all the involved processes keep waiting for each other to get completed. | 2) In this situation, the high priority processes are executed, and low priority processes are blocked. |
| 3) In a deadlock, none of them can execute because they are waiting for the other process to complete. | 4) In starvation, the low priority process starves due to lack of resources. |
| 5) The resources are blocked by the process. | 5) Resources are utilized by high priority process continuously. |
| 6) Deadlock can be prevented by avoiding conditions such as mutual exclusion, hold and wait, no preemption and circular wait. | Starvation can be prevented using the "Aging" technique. |

Bankers algorithm Example

| PI D | Allocation A B C D | Max A B C D | Available A B C D |
|---------|-----------------------|----------------|----------------------|
| P0 | 2 0 0 1 | 4 2 1 2 | 3 3 2 1 |
| P1 | 3 1 2 1 | 5 2 5 2 | |
| P2 | 2 1 0 2 | 2 3 1 6 | |
| P3 | 1 3 1 2 | 1 4 2 4 | |
| P4 | 1 4 3 2 | 3 6 6 5 | |

- 1) What is Need Matrix?
- 2) Is the system is safe state? Is safe find the safe sequence
- 3) If request from p1 arrives for(1,0,0,0) can request be immediately granted?
- 4) If request from P4 arrives for (0,0,2,0) can it be immediately granted?