

Unit-5

Naming

By Prashant Gautam

BCA 6th Semester

Outlines

- Name, Identifiers and Address
- Structured Naming
- Attribute-Based Naming
- Case Study: The Global Name service

Names, Identifiers, and Addresses

- A name in a distributed system is a string of bits or characters that is used to refer to an entity.
- An entity in a distributed system can be practically anything.
 - Typical examples include resources such as hosts, printers, disks, and files.
 - Other well-known examples of entities that are often explicitly named are processes, users, mailboxes, newsgroups, Web pages, graphical windows, messages, network connections, and so on.

For example, a resource such as a printer offers an interface containing operations for printing a document, requesting the status of a print job, and the like. Furthermore, an entity such as a network connection may provide operations for sending and receiving data, setting quality-of-service parameters, requesting the status, and so forth.

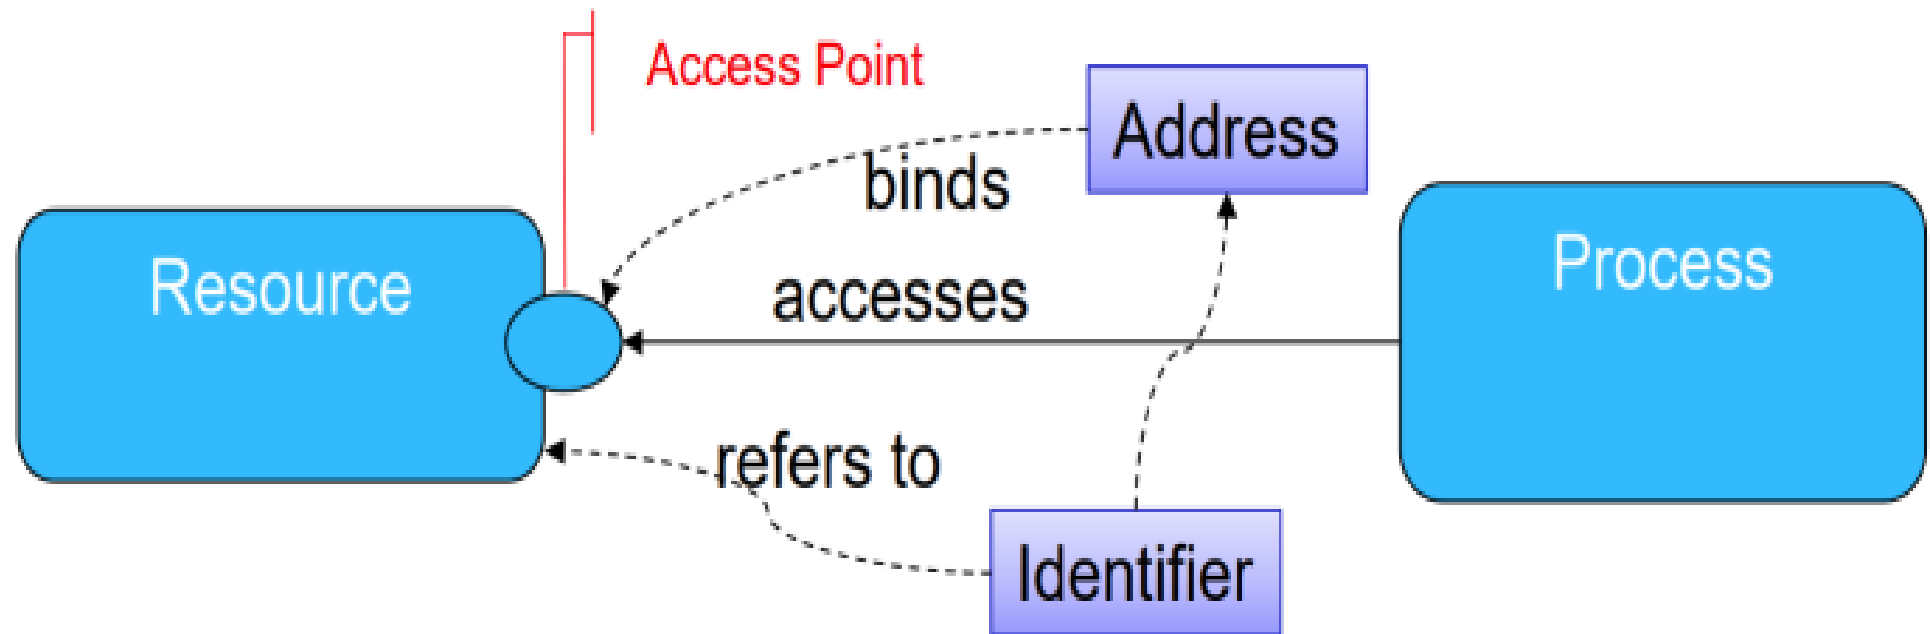
- To operate on an entity, it is necessary to access it, for which we need an access point.
- An access point is yet another, but special, kind of entity in a distributed system.
- The name of an access point is called an address.
- The address of an access point of an entity is also simply called an address of that entity.
- An entity can offer more than one access point.

- As a comparison, a telephone can be viewed as an access point of a person, whereas the telephone number corresponds to an address. Indeed, many people nowadays have several telephone numbers, each number corresponding to a point where they can be reached.

In a distributed system, a typical example of an access point is a host running a specific server, with its address formed by the combination of, for example, an IP address and port number (i.e., the server's transport-level address).

- An entity may change its access points in the course of time.
- For example: when a mobile computer moves to another location, it is often assigned a different IP address than the one it had before. Likewise, when a person moves to another city or country, it is often necessary to change telephone numbers as well.
- In a similar fashion, changing jobs or Internet Service Providers, means changing your e-mail address.

- Identifier: a name that uniquely identifies an entity
 - the identifier is unique and refers to only one entity
- Address: the name of an access point, the location of an entity



3 Classes of naming systems

1. FLAT NAMING / UNSTRUCTURED NAMING
2. STRUCTURED NAMING
3. ATTRIBUTE-BASED NAMING

FLAT NAMING / UNSTRUCTURED NAMING

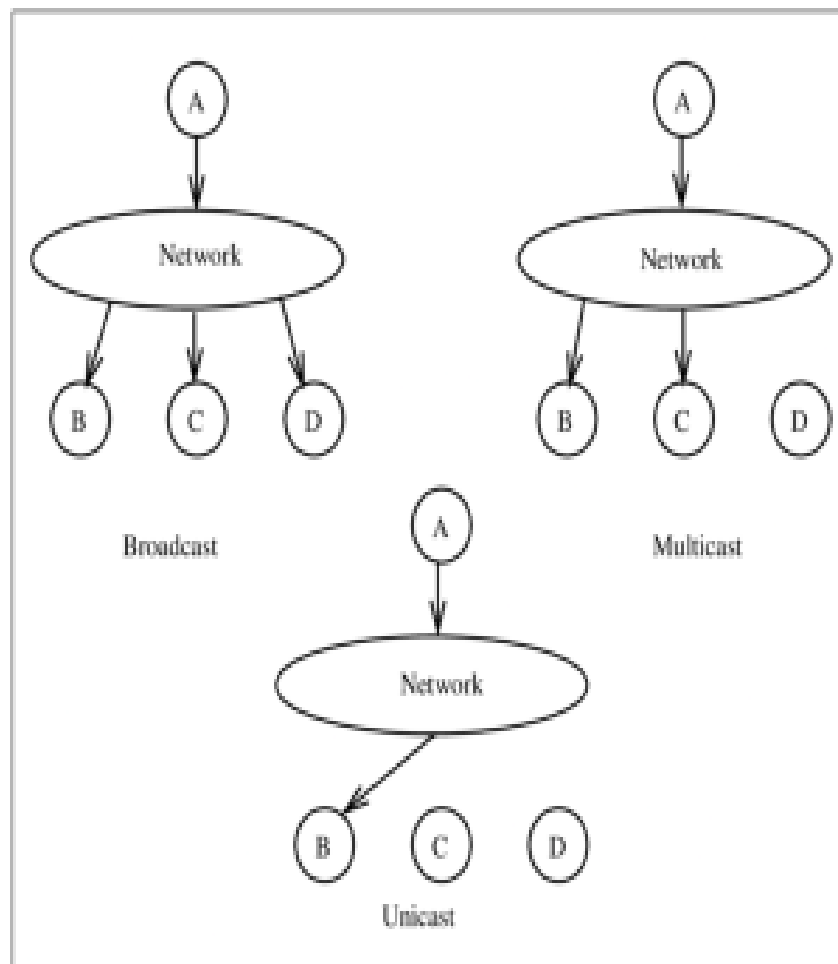
- Identifiers are convenient to uniquely represent entities.
- In many cases, identifiers are simply random bit strings. which we conveniently refer to as unstructured, or flat names.
- An important property of such a name is that it does not contain any information whatsoever on how to locate the access point of its associated entity.
- We will study these aspects:
 - How flat names can be resolved ?
 - How we can locate an entity when given only its identifier ?

FLAT NAMING / UNSTRUCTURED NAMING

- Simple Solution for locating an entity
 - Broadcasting and Multicasting
 - Forwarding Pointers
- Both solutions are applicable only to local-area networks.
- But they often do the job well, making their simplicity particularly attractive.

Broadcasting and Multicasting

BROADCAST VERSUS MULTICAST



BROADCAST	MULTICAST
A method of transferring a message to all recipients simultaneously	A group communication where data transmission is addressed to a group of destination computers simultaneously
Packets are transmitted to all the connected devices in the network	Packets are transmitted to some of the devices in the network
There is no need for group management	Requires group management
Less secure	More secure
More traffic	Less traffic
Slower	Faster
	Visit www.PEDIAA.com

- Consider a distributed system built on a computer network: that offers efficient broadcasting facilities.
- Typically, such facilities are offered by local-area networks and LAN wireless in which all machines are connected to a single cable or equivalent.
- Locating an entity in such an environment is simple: a message containing the identifier of the entity is broadcast to each machine and each machine is requested to check whether it has that entity.
- Only the machines that can offer an access point for the entity send a reply message containing the address of that access point.

Example:

- a machine broadcasts a packet on the local network asking who is the owner of a given IP address.
- When the message arrives at a machine, the receiver checks whether it should listen to the requested IP address.
- If so, it sends a reply packet containing, for example, its Ethernet address.

Problems with Broadcasting:

- Broadcasting becomes inefficient when the network grows.
- Not only is network bandwidth wasted by request messages, but, more seriously, too many hosts maybe interrupted by requests they cannot answer.
- One possible solution is to switch to multicasting, by which only a restricted group of hosts receives the request.
- For example, Ethernet networks support data-link level multicasting directly in hardware.

- Multicasting can also be used to locate entities in point-to-point networks.
- For example, the Internet supports network-level multicasting by allowing hosts to join a specific multicast group.
- Such groups are identified by a multicast address.
- When a host sends a message to a multicast address, the network layer provides a best-effort service to deliver that message to all group members.

- A multicast address can be used as a general location service for multiple entities.
- For example, consider an organization where each employee has his or her own mobile computer.
- When such a computer connects to the locally available network, it is dynamically assigned an IP address.
- In addition, it joins a specific multicast group. When a process wants to locate computer A, it sends a "where is A?" request to the multicast group.
- If A is connected, it responds with its current IP address.

Assignment-VI (Deadline: Feb 28)

- Explain the merits and demerits of broadcasting and multiple casting in locating entities.

Forwarding Pointers

- The principle is simple: when an entity moves from A to B, it leaves behind in A a reference to its new location at B.
- Advantages:
 - simplicity: as soon as an entity has been located, for example by using a traditional naming service, a client can look up the current address by following the chain of forwarding pointer
- Drawbacks
 - a chain for a highly mobile entity can become so long that locating that entity is prohibitively expensive.
 - the vulnerability to broken links

Explore more on Forward pointers

STRUCTURED NAMING

- Flat names:
 - good for machines, but are generally not very convenient for humans to use.
- structured names:
 - composed from simple, human-readable names.

In this section, we concentrate on structured names and the way that these names are resolved to addresses.

Name Spaces

- Names are commonly organized into what is called a name space.
- Name spaces for structured names can be represented as a labeled, directed graph with two types of nodes.

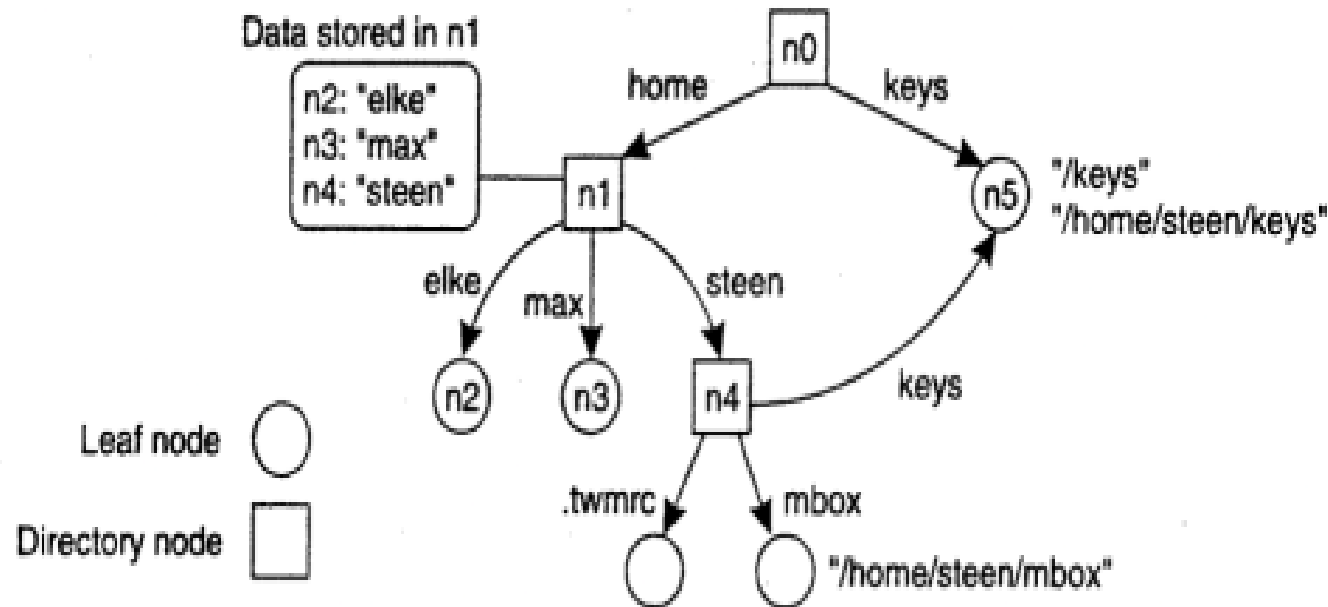


Figure 5-9. A general naming graph with a single root node.

Name Spaces

- A **leaf node** represents a named entity and has the property that it has no outgoing edges.
- A **leaf node** generally stores information on the entity it is representing—for example, its address—so that a client can access it.
- Alternatively, it can store the state of that entity, such as in the case of file systems in which a leaf node actually contains the complete file it is representing.
- A **directory node** has a number of outgoing edges, each labeled with a name.
- A **directory node** stores a table in which an outgoing edge is represented as a pair (edge label, node identifier). Such a table is called a directory table.

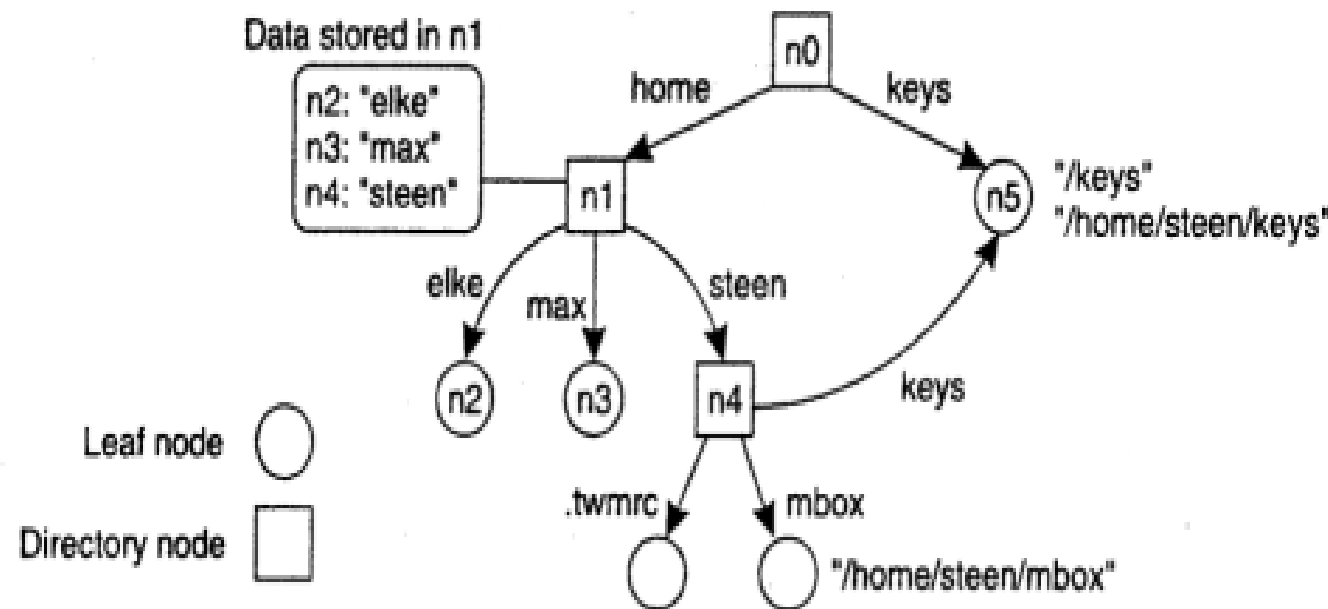


Figure 5-9. A general naming graph with a single root node.

Name Resolution

- Name spaces offer a convenient mechanism for storing and retrieving information about entities by means of names.
- More generally, given a path name, it should be possible to look up any information stored in the node referred to by that name.
- The process of looking up a name is called **name resolution**.

DNS

- We use to remember “human-readable” machine name
 - we have the name hierarchy
 - E.g., www.facebook.com
- But machines in Internet use IP address
 - E.g., 31.13.84.33
 - Application communication use IP addresses and ports
- DNS
 - Mapping from the domain name hierarchy to IP addresses

Information in records of DNS namespace

Type of record	Associated entity	Description
SOA	Zone	Holds information on the represented zone
A	Host	Contains an IP address of the host this node represents
MX	Domain	Refers to a mail server to handle mail addressed to this node
SRV	Domain	Refers to a server handling a specific service
NS	Zone	Refers to a name server that implements the represented zone
CNAME	Node	Symbolic link with the primary name of the represented node
PTR	Host	Contains the canonical name of a host
HINFO	Host	Holds information on the host this node represents
TXT	Any kind	Contains any entity-specific information considered useful

Source: Andrew S. Tanenbaum and Maarten van Steen, Distributed Systems – Principles and Paradigms, 2nd Edition, 2007, Prentice-Hall

ATTRIBUTE-BASED NAMING

- A tuple (attribute, value) can be used to describe a property
 - E.g., („country“, „Austria“), („language“, „German“)
- A set of tuples (attribute, value) can be used to describe an entity

AustriaInfo

Attribute	Value
CountryName	Austria
Language	German
MemberofEU	Yes
Capital	Vienna

- Employ (attribute, value) tuples for describing entities
 - Why flat and structured naming are not enough?
- Also called directory services.
- Name Resolution
 - Usually based on querying mechanism
 - Querying usually deal with the whole space
- Implementation
 - LDAP (Lightweight Directory Access Protocol)
 - RDF (Resource Description Framework): Framework for semantics

Assignment-VII (Deadline: Mar 3)

- Explore more about LDAP, how it works, considering real world example.
- Explore more about RDF, how it works, considering real world example.